Dubai Government-Owned Sovereign AI Media Operations Platform

Proposal for Dubai Police

A sovereign AI platform proposal centered on self-hosted GLM-5.2, governed model routing, human approval, and Dubai government-controlled data residency.

Prepared For	Dubai Police
Document Type	Proposal
Date	27 June 2026
Core Idea	Government-owned sovereign AI foundation with no operational data leaving approved Dubai government servers.

Key Proposal Points

- Self-hosted GLM-5.2 as primary open-weight model candidate.
- Hybrid model catalog for open-source and approved sovereign closed-source models.
- Four product agents: Media Release, Campaign Creative, Screen Adaptation, Monitoring Intelligence.
- Mandatory human review and approval before external publication.
- Reusable AI foundation for future Dubai government use cases.

Executive Summary

This plan defines a sovereign AI Media Operations Platform for Dubai Police. The platform will run within Dubai government-controlled infrastructure and will keep all operational data, prompts, media assets, embeddings, logs, and generated outputs inside approved Dubai government servers.

The platform will use GLM-5.2 as the primary open-weight foundation model candidate for Arabic and English generation, reasoning, summarization, translation, and agentic workflows. The system will also support a governed catalog of additional open-source and approved closed-source models, selected per task through a secure Model Gateway.

The objective is to build more than a single media application. This will become a reusable Dubai government AI foundation that can support media operations first, then expand into future government use cases such as knowledge assistants, policy analysis, multilingual document intelligence, public-service communication, and citizen-service automation.

Platform Vision

Dubai Police will own and operate a secure AI platform that can:

- Generate Arabic and English media releases.
- Create campaign concepts, slogans, and multilingual content.
- Adapt approved creatives for digital screens.
- Monitor approved media sources and produce intelligence summaries.
- Enforce human review and approval before publication.
- Keep all sensitive information inside Dubai government infrastructure.
- Route workloads across multiple approved models.
- Maintain full traceability of prompts, models, outputs, reviewers, and approvals.
- Scale into a shared AI operating layer for wider Dubai government use.

Guiding Principles

- Data sovereignty by design.
- No live operational data leaves approved Dubai government infrastructure.
- Open-weight models are preferred where they meet quality and security needs.
- Closed-source models are allowed only if hosted inside a sovereign, approved environment.
- Human approval is mandatory before any external publication.
- Safety and governance controls are enforced by the platform, not only by model prompts.
- All prompts, models, source documents, outputs, and approvals are auditable.
- The architecture must support many models and future use cases.

- Arabic-first quality must be measured using Dubai Police content and UAE-specific review criteria.

Security Note: External public SaaS model APIs must not be used for live Dubai Police content unless a formally approved sovereign deployment guarantees no data leaves approved Dubai government infrastructure. The secure approach is self-hosted open-weight models or closed-source models deployed inside the approved sovereign environment with private networking, no provider training, no prompt retention, and audited controls.

GLM-5.2 Model Strategy

GLM-5.2 is selected as the primary open-weight foundation model candidate for the first platform release.

Rationale:

- Open-weight deployment supports local hosting.
- MIT licensing supports broad internal government use.
- Strong long-context capability supports lengthy policies, historic media content, campaign libraries, audit records, and monitoring reports.
- Strong reasoning and agentic capability supports structured workflows and tool use.
- Arabic and English workflows can run without sending sensitive data to external AI providers.
- It can act as the main model behind a government-owned Model Gateway.

Project premise:

- GLM-5.2 is positioned as a top Arabic-capable open-source model and is claimed to rank number 2 for Arabic language performance.
- The project must validate this claim during PoC using Dubai Police Arabic content, Modern Standard Arabic samples, UAE communication examples, translation tests, summarization tests, and policy-compliance tests.
- Production approval must be based on measured benchmark results, reviewer feedback, latency, cost, and security fit.

Hybrid Model Catalog

The platform will not depend on a single model. It will use a governed model catalog managed by Dubai Police or the approved Dubai government AI operations team.

Model categories:

- Foundation LLMs.
- Arabic-specialized LLMs.
- Translation models.
- Embedding models.

- Vision-language models.
- OCR models.
- Content safety models.
- Sentiment and classification models.
- Speech-to-text models for future audio/video workflows.
- Text-to-speech models for future voice workflows.

Each model record must include:

- Model name.
- Version.
- License.
- Hosting location.
- Hardware profile.
- Supported languages.
- Supported modalities.
- Approved use cases.
- Restricted use cases.
- Benchmark scores.
- Safety evaluation results.
- Known limitations.
- Owner.
- Approval date.
- Review or retirement date.

Model Gateway

The Model Gateway is the control layer for all AI inference. Applications and agents must call the Model Gateway, not model servers directly.

Responsibilities:

- Route requests to approved models.
- Enforce data classification.
- Enforce prompt templates.
- Apply policy and safety checks.
- Block unapproved external endpoints.
- Track model version and prompt version.
- Track token usage, latency, throughput, and infrastructure cost.
- Support fallback across approved models.

- Support OpenAI-compatible APIs for easier integration.
- Maintain complete audit records.

Routing examples:

Workload	Preferred Model Path
Arabic press release generation	GLM-5.2 or validated Arabic-specialized local model
English adaptation	GLM-5.2 or local translation-specialized model
Campaign ideas and slogans	GLM-5.2 with brand-guideline retrieval
Long policy summarization	GLM-5.2 long-context mode
Monitoring report synthesis	GLM-5.2 with retrieval and classification models
Sentiment analysis	Smaller local classifier if quality is sufficient
Entity extraction	Smaller local NLP model or GLM-5.2
Screen layout analysis	Local vision-language model
OCR	Local OCR model
Safety moderation	Local safety classifier plus rule-based policy engine
High-risk communication	Strongest approved model plus senior approval workflow

Product Modules

1. Media Release Agent

Purpose:

- Generate and refine official media releases in Arabic and English.

Capabilities:

- Read brief.
- Generate Arabic draft.
- Create English adaptation.
- Translate between Arabic and English.
- Summarize background material.
- Improve tone and clarity.
- Prepare captions and hashtags.
- Assemble media package.
- Align with approved terminology and tone-of-voice.
- Route output for review and approval.

Primary model path:

- GLM-5.2 with retrieval from approved media guidelines, historical releases, prohibited wording, and terminology glossary.

2. Campaign Creative Agent

Purpose:

- Generate campaign concepts, slogans, messaging, and social content.

Capabilities:

- Read campaign brief.
- Generate creative concepts.
- Generate campaign slogans.
- Produce Arabic and English social media content.
- Recommend approved asset bundles.
- Create channel-specific content variants.
- Align with brand guidelines.
- Route output for review and approval.

Primary model path:

- GLM-5.2 for concepts, slogans, and structured campaign planning.
- Local vision or asset-analysis model where visual review is required.

3. Screen Adaptation Agent

Purpose:

- Adapt approved creative content for Dubai Police digital screens.

Capabilities:

- Accept approved creative.
- Match screen specifications.
- Check aspect ratio.
- Check safe zones.
- Check readability.
- Validate layout and resolution.
- Recommend resize/adaptation actions.

- Generate preview sheets.
- Flag risks before publication.

Primary model path:

- Local vision-language model for visual analysis.
- GLM-5.2 for reasoning, explanation, and adaptation instructions.

4. Monitoring Intelligence Agent

Purpose:

- Monitor approved sources and generate intelligence summaries, alerts, and reputation reports.

Capabilities:

- Scan approved sources.
- Classify content.
- Detect sentiment.
- Detect topics.
- Track entities.
- Flag misinformation.
- Identify fraud/risk indicators.
- Generate alert reports.
- Generate executive summaries.
- Preserve source references.

Primary model path:

- GLM-5.2 for synthesis and executive summaries.
- Local classifiers for sentiment, topic detection, entity extraction, fraud signals, and misinformation indicators.

Knowledge and Retrieval Layer

The platform will use a private retrieval-augmented generation layer so model outputs are grounded in approved sources.

Knowledge sources:

- Brand guidelines.
- Tone-of-voice documents.
- Approved Arabic and English press releases.

- Historical campaign samples.
- Legal and policy documents.
- Prohibited terms.
- Approved terminology glossary.
- Crisis communication playbooks.
- Digital screen specifications.
- Trusted monitoring sources.
- Internal communication standards.

Core components:

- Document ingestion pipeline.
- OCR and parsing.
- Metadata tagging.
- Embedding generation inside government infrastructure.
- Vector database.
- Search index.
- Source citation and evidence linking.
- Access-controlled retrieval.
- Data retention controls.

Human Review and Approval

No AI-generated output can be published automatically.

Workflow:

1. User submits request.
2. Request is classified.
3. Model Gateway selects approved model.
4. Agent generates draft using approved prompt and retrieved evidence.
5. Safety and policy checks run.
6. Reviewer validates the output.
7. Approver approves, rejects, or requests edits.
8. Approved content is exported, distributed, or published through approved channels.
9. Every action is logged.

High-risk outputs require additional review:

- Public safety communications.
- Crisis communications.

- Legal or regulatory statements.
- Reputation-sensitive content.
- High-visibility public messaging.
- Content flagged for misinformation, bias, unsupported claims, or sensitive wording.

Governance Framework

Required governance controls:

- Model catalog governance.
- Prompt versioning.
- Prompt approval workflow.
- Model approval workflow.
- Human review workflow.
- Approval history.
- Policy validation.
- Content safety validation.
- Source citation requirements.
- Sensitive-topic routing.
- Red-team testing.
- Audit logging.
- Role-based access control.
- Environment segregation.
- Department or tenant segregation if required.

Governance records must capture:

- User request.
- Source documents retrieved.
- Agent invoked.
- Model used.
- Model version.
- Prompt used.
- Prompt version.
- Generated output.
- Safety results.
- Policy results.
- Reviewer decision.
- Approver decision.
- Final output.

- Publication or export action.

Security and Data Residency

The platform must enforce data residency at every layer.

Required controls:

- Private network access for AI workloads.
- No default public internet egress from inference services.
- Egress allow-listing through approved security gateways.
- No prompt or document upload to public AI APIs.
- Encryption in transit.
- Encryption at rest.
- Hardware-backed key management or approved government key vault.
- Secrets management.
- Private model registry.
- Private container registry.
- Private artifact storage.
- Central SIEM integration.
- Immutable logging for critical approval events.
- Data classification labels.
- Retention and deletion controls.
- Backup and disaster recovery.

Application security controls:

- SSO.
- MFA.
- RBAC.
- Least-privilege access.
- API authentication.
- API authorization.
- Input validation.
- Dependency vulnerability management.
- Secure coding practices.
- Audit trail retention.

AI security controls:

- Prompt-injection detection.

- Retrieval sanitization.
- Tool-use restrictions.
- Model output validation.
- Content safety classification.
- Misinformation risk checks.
- Human approval checkpoints.
- Red-team testing before go-live.

Infrastructure Architecture

The platform requires a GPU-backed sovereign AI environment.

Core infrastructure:

- Dubai government data center, sovereign cloud, or approved private cloud.
- Kubernetes or OpenShift cluster.
- GPU node pools for large models.
- Smaller GPU node pools for embeddings, OCR, classification, and safety models.
- High-throughput encrypted storage for model weights.
- Private container registry.
- Private model artifact registry.
- Relational database for workflows and metadata.
- Object storage for media assets.
- Vector database for retrieval.
- API gateway.
- Load balancer.
- Central logging.
- Observability stack.
- SIEM integration.
- Backup and disaster recovery environment.

Model serving layer:

- vLLM, SGLang, or another approved serving framework.
- OpenAI-compatible inference endpoints.
- Batch inference queue.
- Streaming inference support.
- Model health checks.
- GPU utilization monitoring.
- Latency and throughput monitoring.

- Cost and capacity dashboards.

Important sizing note:

- GLM-5.2 is a large model and requires serious GPU capacity.
- The project must benchmark full-size, FP8, quantized, and smaller-model options.
- Smaller specialist models should handle routine tasks where they meet accuracy requirements.
- GPU sizing must be finalized after the PoC and expected workload analysis.

Benefits

1. Data Sovereignty

All sensitive prompts, documents, source material, generated outputs, and logs remain inside Dubai government-controlled infrastructure.

Benefits:

- Stronger confidentiality.
- Better compliance alignment.
- Full control over retention and deletion.
- Reduced exposure to foreign cloud AI providers.
- Clear ownership of operational data.

2. Government AI Ownership

Dubai government owns the AI operating layer, model gateway, governance framework, and reusable infrastructure.

Benefits:

- Ability to scale beyond media operations.
- Government-controlled roadmap.
- Reusable AI services for future departments and use cases.
- Less dependency on a single commercial model provider.

3. Arabic-First Quality

Arabic quality will be measured and governed locally.

Benefits:

- Dubai Police terminology can be standardized.
- UAE communication style can be reviewed by local experts.
- Modern Standard Arabic, bilingual, and UAE-specific examples can be benchmarked.
- Prompt and retrieval improvements can be made continuously.

4. Hybrid Model Flexibility

The platform can use multiple open-source and approved closed-source models.

Benefits:

- Best model can be selected per task.
- Models can be upgraded without rebuilding the application.
- Routine tasks can use smaller efficient models.
- High-risk tasks can use stronger models and stricter approval.
- Model vendor changes do not break the platform.

5. Long-Term Cost Control

Initial infrastructure cost is higher, but high-volume usage can be controlled internally.

Benefits:

- Reduced per-token public API dependency.
- Shared GPU infrastructure across multiple use cases.
- Smaller models can reduce routine workload cost.
- Capacity planning can be aligned with government demand.

6. Stronger Auditability

Self-managed AI infrastructure allows deeper governance evidence.

Benefits:

- Complete prompt and model traceability.
- Complete source and retrieval traceability.
- Clear review and approval history.
- Better incident investigation.
- Stronger compliance reporting.

7. Future Expansion

The same AI foundation can later support:

- Internal knowledge assistants.
- Policy summarization.
- Legal drafting support.
- Public-service communication.
- Citizen-service automation.
- Call-center augmentation.
- Incident reporting.
- Multilingual document intelligence.
- Training and onboarding assistants.
- Department-specific AI agents.

Implementation Roadmap

Phase 0: Sovereign AI Strategy and Model Validation

Duration: Week 1-3

Key activities:

- Confirm data residency rules.
- Confirm approved hosting environment.
- Confirm Dubai government infrastructure ownership model.
- Validate GLM-5.2 license and deployment options.
- Validate Arabic performance claim through benchmark design.
- Define Arabic and English benchmark dataset.
- Select secondary model candidates.
- Define model evaluation criteria.
- Define security and egress policy.
- Define AI governance operating model.

Milestones:

- Sovereign AI strategy approved.
- Data-residency policy confirmed.
- Model evaluation framework approved.
- GLM-5.2 PoC approved.
- Benchmark dataset approved.

Phase 1: Infrastructure Sizing and Platform Design

Duration: Week 4-7

Key activities:

- Size GPU infrastructure for GLM-5.2.
- Size smaller-model infrastructure.
- Compare full-size, FP8, quantized, and smaller-model deployments.
- Design Kubernetes/OpenShift platform.
- Design network isolation.
- Design storage, registry, backup, and disaster recovery.
- Design Model Gateway.

- Design observability and cost monitoring.
- Prepare procurement bill of materials.

Milestones:

- AI infrastructure architecture approved.
- GPU sizing completed.
- Procurement BOM completed.
- Model Gateway design approved.
- Security architecture approved.

Phase 2: Sovereign AI Platform Build

Duration: Week 8-15

Key activities:

- Provision GPU cluster.
- Set up Kubernetes/OpenShift.
- Set up private container registry.
- Set up private model registry.
- Deploy model-serving framework.
- Deploy GLM-5.2 PoC endpoint.
- Deploy Model Gateway.
- Set up vector database and search.
- Set up prompt registry.
- Set up audit logging.
- Set up observability.
- Implement network isolation and egress controls.

Milestones:

- GPU cluster operational.
- GLM-5.2 served locally.
- Model Gateway live in staging.
- Prompt registry live.
- Audit and observability live.
- No-external-data-flow control validated.

Phase 3: Application and Agent Build

Duration: Week 12-24

Key activities:

- Build React portal.
- Build backend microservices.
- Implement request intake.
- Implement Media Release Agent.
- Implement Campaign Creative Agent.
- Implement Screen Adaptation Agent.
- Implement Monitoring Intelligence Agent.
- Integrate agents with Model Gateway.
- Integrate retrieval layer.
- Implement approval workflow.
- Implement Governance Dashboard.
- Implement notifications.
- Implement RBAC and SSO.
- Implement audit views and reporting.

Milestones:

- Portal live in staging.
- Four agents implemented.
- Approval workflow implemented.
- Governance Dashboard implemented.
- RBAC and SSO configured.
- Retrieval-grounded generation working.

Phase 4: Evaluation, Security Testing, and UAT

Duration: Week 20-28

Key activities:

- Run Arabic and English quality evaluation.
- Test 97%+ accuracy framework.
- Test hallucination reduction.
- Test source citation workflows.
- Test sensitive-topic escalation.
- Run model red-team testing.
- Run prompt-injection testing.

- Run application security testing.
- Run data-egress testing.
- Run performance and load testing.
- Conduct UAT with Dubai Police users.
- Resolve defects.

Milestones:

- Arabic benchmark report completed.
- Security test report completed.
- Data-residency validation completed.
- UAT completed.
- Defect log cleared.
- Acceptance report submitted.

Phase 5: Training, Handover, and Operational Readiness

Duration: Week 29-31

Key activities:

- Train communications users.
- Train reviewers and approvers.
- Train administrators.
- Train AI operations team.
- Train infrastructure and security teams.
- Deliver user manuals.
- Deliver admin guide.
- Deliver prompt governance guide.
- Deliver model operations runbook.
- Deliver deployment and support documentation.
- Conduct operational readiness review.

Milestones:

- User groups trained.
- AI operations team trained.
- Documentation delivered.
- Handover pack accepted.
- Go-live readiness approved.

Phase 6: Production Go-Live and Hypercare

Duration: Week 32-36 plus ongoing

Key activities:

- Deploy production platform.
- Validate production data-residency controls.
- Monitor model quality.
- Monitor GPU utilization.
- Monitor latency and throughput.
- Monitor review and approval rates.
- Tune prompts.
- Tune routing rules.
- Provide hypercare support.
- Transition to SLA-based operations.

Milestones:

- Platform live in production.
- Go-live report submitted.
- Hypercare completed.
- SLA support activated.
- Optimization backlog created.

Timeline Summary

Phase	Duration	Calendar Weeks
Phase 0 - Sovereign AI Strategy and Model Validation	3 weeks	Week 1-3
Phase 1 - Infrastructure Sizing and Platform Design	4 weeks	Week 4-7
Phase 2 - Sovereign AI Platform Build	8 weeks	Week 8-15
Phase 3 - Application and Agent Build	13 weeks	Week 12-24
Phase 4 - Evaluation, Security Testing, and UAT	9 weeks	Week 20-28
Phase 5 - Training, Handover, and Readiness	3 weeks	Week 29-31
Phase 6 - Go-Live and Hypercare	5 weeks plus ongoing	Week 32-36 plus ongoing

Overall estimate:

- 36 weeks to production.
- 4 weeks of hypercare after go-live.
- Timeline can reduce if approved GPU infrastructure, Kubernetes platform, model registry, and AI operations tooling already exist.

Workstreams

Workstream A: Sovereign Infrastructure

Deliverables:

- GPU cluster.
- Kubernetes/OpenShift platform.
- Private registry.
- Model registry.
- Inference-serving stack.
- Storage and backup.
- Network security.
- Egress controls.
- Disaster recovery plan.

Workstream B: Model and AI Platform

Deliverables:

- GLM-5.2 deployment.
- Model Gateway.
- Model catalog.
- Prompt registry.
- Retrieval layer.
- Embedding service.
- Safety models.
- Benchmark framework.
- Model monitoring.

Workstream C: Application Product

Deliverables:

- Web portal.
- Backend APIs.
- Four agents.
- Approval workflows.
- Governance Dashboard.

- Notifications.
- Reporting.
- Admin screens.

Workstream D: Governance and Security

Deliverables:

- AI governance policy.
- Prompt governance process.
- Model approval process.
- Audit framework.
- Role matrix.
- Sensitive-topic policy.
- Red-team test plan.
- Security test reports.
- Data-residency evidence.

Workstream E: Adoption and Operations

Deliverables:

- Training.
- Documentation.
- Support runbook.
- AI operations runbook.
- Infrastructure runbook.
- SLA model.
- Knowledge transfer.
- Hypercare plan.

Team Requirements

Vendor roles:

- Project Manager.
- Business Analyst.
- Solution Architect.
- Sovereign AI Architect.
- AI Infrastructure Architect.
- MLOps Engineer.
- GPU Infrastructure Engineer.
- AI/ML Engineers.
- Prompt Engineer.
- Backend Developers.
- Frontend Developers.
- Database Engineer.
- DevOps Engineer.
- Security Engineer.
- QA/Test Engineers.
- Arabic Language QA Specialists.
- UX/UI Designer.
- Technical Writer.

Dubai Police / Dubai Government roles:

- Project Sponsor.
- Product Owner.
- Media Operations Representatives.
- Governance Representatives.
- IT Infrastructure Representatives.
- Cybersecurity Representatives.
- Identity Management Representative.
- Data Protection / Compliance Representative.
- AI Operations Owner.
- Model Governance Owner.
- UAT Reviewers and Approvers.

Deliverables

Application deliverables:

- Web portal.
- Media Release Agent.
- Campaign Creative Agent.
- Screen Adaptation Agent.
- Monitoring Intelligence Agent.
- Governance Dashboard.
- Approval workflow engine.
- Audit and reporting.
- Admin module.
- Notification module.

Sovereign AI deliverables:

- Model Gateway.
- GLM-5.2 deployment.
- Model catalog.
- Prompt registry.
- Benchmark framework.
- Retrieval layer.
- Vector database.
- Safety and policy validation layer.
- Model monitoring dashboard.
- AI operations runbook.

Infrastructure deliverables:

- GPU serving environment.
- Kubernetes/OpenShift deployment.
- Private model registry.
- Private container registry.
- Secure network design.
- Observability and SIEM integration.
- Backup and disaster recovery design.

Governance deliverables:

- AI governance framework.

- Data-residency evidence report.
- Model approval policy.
- Prompt approval policy.
- Red-team report.
- Security test report.
- Arabic benchmark report.
- Human approval process.

Acceptance Criteria

The platform is accepted when:

- No production prompt, source document, media file, embedding, log, or generated output leaves approved Dubai government infrastructure.
- GLM-5.2 or the selected primary model passes agreed Arabic and English quality benchmarks.
- The four agents meet approved functional requirements.
- Human approval is mandatory before publication.
- Prompt and model versions are logged.
- Output source references are retained where required.
- Sensitive-topic escalation works.
- RBAC works at application and API layers.
- Environment segregation works.
- Department or tenant segregation works if required.
- Security testing passes.
- Data-egress testing passes.
- UAT is signed off.
- Documentation and runbooks are delivered.

Risks and Mitigation

Risk	Impact	Mitigation
GLM-5.2 Arabic performance does not meet expectations	Quality and acceptance risk	Benchmark early with Dubai Police content and keep secondary model candidates ready.
GPU procurement delays	Timeline risk	Start procurement during Phase 0 and identify interim approved GPU capacity.
Infrastructure sizing is underestimated	Performance and cost risk	Run load tests and size for peak workflows.
Closed-source model use violates data residency	Compliance risk	Allow closed models only through approved sovereign deployment.
Open model safety gaps	Security and reputational risk	Use safety classifiers, policy engine, red-team testing, and human approval.
Model hallucination	Public communication risk	Use retrieval grounding, source references, automated validation, and reviewer approval.
Prompt injection through monitoring sources	Security risk	Sanitize retrieved content, isolate tools, apply prompt-injection filters, and restrict agent actions.
Operational complexity	Support risk	Create AI operations runbooks, dashboards, and dedicated operations training.
Higher initial infrastructure cost	Budget risk	Position infrastructure as reusable government AI platform supporting many future use cases.

Commercial Implications

Cost categories:

- One-time product implementation.
- GPU and data-center infrastructure.
- Model deployment and MLOps.
- Application development.
- Security and compliance testing.
- Training and handover.
- Hypercare and SLA support.
- Ongoing AI operations.
- Hardware refresh and capacity expansion.

Commercial benefits:

- Reusable government AI infrastructure.
- Reduced external AI API dependency.
- Better long-term control over usage cost.
- Shared GPU platform for future departments.
- Stronger ownership of model operations and data.

Open Decisions

The following decisions are required before final design:

- Exact Dubai government hosting environment.
- Hardware procurement or sovereign GPU cloud capacity.
- Final GLM-5.2 deployment profile: full precision, FP8, quantized, or smaller derivative.
- Approved secondary models.
- Whether any closed-source models can be deployed sovereignly.
- Final benchmark suite and acceptance thresholds.
- Department or tenant segregation model.
- Approved social monitoring and publishing integrations.
- Retention period for prompts, outputs, embeddings, and audit logs.
- Disaster recovery targets: RPO and RTO.
- Production SLA and support hours.

Immediate Next Steps

1. Approve the sovereign AI direction.
2. Run a GLM-5.2 Arabic and English benchmark PoC.
3. Confirm Dubai government data-residency and hosting constraints.
4. Size GPU infrastructure for the expected workload.
5. Define model catalog governance.
6. Finalize procurement and platform architecture.
7. Begin parallel work on the application portal and sovereign AI platform.

Reference Notes

- GLM-5.2 model card states MIT licensing, local serving options, and 1M-token context support: https://huggingface.co/zai-org/GLM-5.2
- vLLM documentation provides GLM-5.2 serving guidance and hardware assumptions for FP8 deployment: https://recipes.vllm.ai/zai-org/GLM-5.2
- Z.ai GLM-5.2 announcement provides additional model positioning and release context: https://z.ai/blog/glm-5.2
- The Arabic number 2 ranking claim should be confirmed through project benchmark evidence before being used in formal commercial or contractual material.